

DEAL SCREEN



Open-source AI models, agents, and decentralized training infrastructure
Austin, TX | Founded 2022-23

Round	New Financing, Jul. 2026 — reported, not yet closed
Size / Pre / Post	\$75M+ at a reported \$1.5Bn valuation; pre/post and instrument structure TBD
Lead Investor	Robot Ventures — reported lead, existing investor
Participants	Union Square Ventures — reported new investor; others TBD
Sourcing	ID8 seeking allocation; no confirmed access path
Step-Up	~2.9x from the prior \$523M equity post-money valuation*
Total Raised	\$70M to date; ~\$145M+ pro forma if the round closes

* Subject to confirmation that the current financing is priced on a comparable equity basis.

July 2026

STRICTLY CONFIDENTIAL

What the Company Does

Nous Research is an open-source AI lab that builds frontier open-weight models, autonomous agents, and decentralized training infrastructure. The company delivers three core products:

- **Hermes models**, a family of open-weight LLMs (latest: Hermes 4, Aug '25) tuned for reasoning, tool use, and steerability with minimal content restrictions; trained on a primarily newly synthesized dataset developed through the DataForge and Atropos pipeline.
- **Hermes Agent**, an MIT-licensed self-improving agent (launched Feb '26) that runs persistently, writes reusable skills from experience, and connects to 20 messaging platforms. MoA 2.0 (Jun '26) blends multiple frontier models (Claude, GPT, DeepSeek) into a single higher-scoring virtual model.
- **Psyche**, a decentralized network that coordinates GPUs across different locations to train large AI models together. Built using Nous's DisTrO technology, it aims to make model training less dependent on costly centralized data centers.

Nous serves developers, researchers, and enterprises building AI applications. Monetization is early: paid Hermes Agent subscriptions (\$20-200/mo) are the current path; no company revenue has been disclosed. A future Psyche token could add a network incentive and value-accrual mechanism, but a token is not company revenue, and its relationship to equity value is undefined. Open releases drive distribution.

Why Consider

- **Breakout Developer Distribution:** Hermes Agent has surpassed OpenClaw in attributed OpenRouter usage, with ~215K GitHub stars, ~40K forks, and 50M+ company-reported Hermes 3 downloads. Its learning loop creates reusable skills that improve with use, supporting strong developer engagement and potentially increasing product stickiness over time.
- **Capital-Efficient Technical Execution:** Nous improves open foundation models through differentiated post-training and synthetic data rather than funding costly models from scratch. Hermes 4 delivered strong reasoning benchmarks, while co-founder Jeffrey Quesnelle's YaRN research has been adopted across leading open-model ecosystems.
- **Multiple Paths to Value Creation:** MoA 2.0 combines outputs from leading frontier models, while DisTrO and Psyche target decentralized model training. Both remain early, but either could become a meaningful platform if Nous proves scalability and monetization.
- **Large Enterprise Opportunity:** Paid Hermes Agent subscriptions provide an early revenue path, but the larger long-term opportunity is to let companies privately deploy Hermes, connect it to internal systems, and customize agents for company-specific workflows while retaining control over their data and model choice. Hermes's open, self-hosted, and model-agnostic architecture is well suited to enterprises prioritizing privacy and reduced vendor lock-in.
- **Strong Backers and Active Re-Pricing:** The reported \$75M+ financing at a \$1.5Bn valuation is being led by Robot Ventures with USV participating. Paradigm led the prior round, and the cap table includes Together AI and prominent founders across the open-AI ecosystem.

Management Team

Jeffrey Quesnelle: Co-Founder & CEO

University of Michigan M.S. in Computer Science. Equal-contribution co-author of YaRN, a long-context method used by Meta and DeepSeek, and co-author of DeMo distributed-training research with Diederik Kingma. Previously led software development at Intrepid Control Systems.

Karan Malhotra: Co-Founder, Head of Behavior

Former ML researcher at the Stanford Brain Stimulation Lab. Leads model behavior, alignment, and strategy across Hermes, WorldSim, and DisTrO, and serves as a primary public voice for the company.

Teknum: Co-Founder, Head of Post-Training

Formerly at Stability AI and creator of the Hermes model line and GPTeacher dataset. Leads post-training, data curation, and HermesBench and is a prominent figure in open-source fine-tuning.

Shivani Mitra: Co-Founder & President

Founding CEO through the Series A, responsible for building the team, brand, and operations and raising approximately \$70M from pre-seed through Series A. Also Managing Partner of emerging-technology investment firm Prometheia.

Cap Table

Seed (Jan 2024) – \$5.2M @ \$22M post (PB)

Distributed Global (co-lead, entered)

OSS Capital (co-lead, entered)

Angels: Balaji Srinivasan, Vipul Ved Prakash (Together AI), Gavin Uberti (Etched), Alexander Atallah (OpenSea / OpenRouter), Chris Prucha (Notion), Sahil Chaudhary (Glaive)

Seed Extension (Jun 2024) – ~\$15M (incl. \$3.55M debt)

Delphi Ventures, North Island Ventures (entered)

Mozilla Ventures, Breed VC, Mechanism Capital, Collab+Currency, Cadenza Capital

Series A / A-1 (Apr 2025) – \$50M @ \$523M post equity; \$1Bn token valuation

Paradigm (lead, entered)

Delphi Digital, Robot Ventures, Faction Ventures (entered)

North Island Ventures, Distributed Global (follow-on); Rajiv Gokal, Together AI (strategic, per CB Insights)

Rumored Round (Jul 2026) – \$75M+ @ \$1.5Bn (not closed)

Robot Ventures (lead, follow-on)

Union Square Ventures (NEW)

T1

Per TechCrunch (Jul 13, 2026), three sources; strong investor interest reported

Financial Snapshot

Metric	Seed ('24)	Series A ('25)	Latest ('26)
Post-Money	\$22M (Jan '24, PB)	\$523M equity / \$1Bn token	\$1.5Bn (rumored, TC)
Revenue	Not disclosed	Not disclosed	Not disclosed; Agent tiers \$20-200/mo live
Total Raised	\$5.2M	\$70M (incl. \$4.5M debt)	\$70M; ~\$145M if rumored round closes
OpenRouter Daily Tokens (attributed)	n/a	n/a (Agent launched Feb '26)	224Bn (May 10) to 458Bn (May 21)
GitHub Stars (Hermes Agent)	n/a	n/a	~215K + ~40K forks (Jul '26)
Model Downloads	n/a	50M+ Hermes 3, all platforms, company-reported	33M+ Hermes-family, Hugging Face only (different basis from the 50M+, not a decline)
Headcount	~8 (est.)	20 (PB, Apr '25)	~48 (LinkedIn est., unverified)

Deal Scoring

Category	Score
Lead / Round Dynamics	2 / 4
Founder / Team Quality	3 / 4
Fundamentals	2 / 4
Return Potential	3 / 4
AI Score	4 / 4
Terms	2 / 4
Overall	2.8 / 4.0 (pending diligence)

Risks & Mitigants

- **Open-Source Monetization:** Nous's models and agent software are widely available for free, creating a risk that usage grows without meaningful revenue capture. Mitigant: Paid hosting, premium agent features, enterprise deployments, and potential network economics provide multiple paths to monetize the developer base.
- **Model-Layer Commoditization:** Free releases from larger, better-funded labs could reduce the differentiation of Hermes models. Mitigant: Nous is expanding beyond standalone models through Hermes Agent, MoA orchestration, proprietary post-training workflows, and its developer community.
- **Intense Agent Competition:** Hermes competes with OpenClaw and well-funded products from OpenAI, Anthropic, and other frontier labs. Mitigant: Hermes has strong developer adoption, reusable skill learning, model-agnostic architecture, and a self-hosted offering that appeals to users seeking flexibility and control.
- **Limited Model Guardrails:** Hermes's minimal content restrictions attract developers but may create regulatory, platform, and enterprise adoption challenges. Mitigant: Nous can add configurable safety and compliance controls at the hosted Agent layer while preserving the openness of its underlying models.
- **Execution Risk Across Multiple Products:** Nous is simultaneously building models, an agent platform, orchestration technology, and decentralized training infrastructure, which could stretch a relatively small team. Mitigant: The products share a common technical stack and developer ecosystem, while prior execution across Hermes, DisTrO, and Psyche demonstrates a strong pace of development.

Developer Sentiment (Reddit / GitHub / HN)

- **Clear Differentiation vs. OpenClaw:** OpenClaw is a gateway first, a switchboard connecting an AI to 50+ messaging apps with a 44K+ library of community skills it downloads but never improves. Hermes is an agent first, a runtime that writes and refines its own skills from experience and remembers across sessions, so it compounds on a user's specific workflows over time (the closest thing to a per-user moat in open-source agents, and why it burns more inference per user than any rival). Nous is pressing that edge with a one-command migration tool (hermes claw migrate) that imports an OpenClaw user's settings, memories, and skills, deliberately raiding the incumbent's installed base.
- Community sentiment is divided. Supporters cite easier setup, stronger memory, and reusable skills, while critics question self-evaluation, stability, and possible promotional activity. Many developers appear to use Hermes and OpenClaw together rather than viewing them as mutually exclusive.
- **Top Recurring Criticism, Self-Evaluation:** The most-upvoted complaint (+107) is that Hermes always grades its own work as successful even when the output is wrong ('It always thinks it did a good job. ALWAYS.'). This attacks the core differentiator: if the agent cannot judge success, the skills it saves may encode errors. Also cited: it overwrites manual edits, and stability claims are premature (11 releases vs. OpenClaw's 137).
- **Security Standing:** OpenClaw disclosed 9 CVEs in 4 days (Mar '26, one CVSS 9.9) and a Koi Security audit found 341 of 2,857 ClawHub skills malicious; Hermes has zero agent-specific CVEs (one MEDIUM in v0.8.0 webhooks) and ships conservative defaults (container hardening, read-only rootfs, pre-execution command scanner). Honest read: zero CVEs partly reflects less exposure, not proven hardening, but Hermes gained 6,400+ GitHub stars in a single day during OpenClaw's security crisis, so the market rewarded the contrast.
- **Developer Ecosystem:** Nearly 300 developers have contributed to Hermes, with almost 600 proposed improvements accepted into the product. This suggests a genuinely active open-source community rather than marketing-driven hype. The main watch item is recurring Reddit suspicion that some promotion may be coordinated.

Comp Set

COMPANY	STAGE	VALUATION	REVENUE	MULTIPLE
CLOSEST PRIVATE COMPS - OPEN-MODEL LABS & INFRASTRUCTURE				
Prime Intellect	Early Stage VC	\$1.0Bn (Jul 2026); ~\$150M raised	~\$100M ann. run rate (TC, co-rep.)	~10x run rate
Mistral AI	Later Stage VC	\$13.66Bn (Sep 2025); \$3.18Bn raised (excl. debt)	~\$200M rev 2025 (Forbes, Apr '26)	~68x trailing (see note)
Hugging Face	Late Stage (mark stale, Aug '23)	\$4.5Bn (Aug 2023); \$395M raised	\$70M 2023 (Sakra); ~\$130M 2024 est.	~35x est. revenue (stale)
Nous Research (ID8 entry)	Early Stage VC	\$1.5Bn reported current round; prior \$523M equity / \$1Bn token	Not disclosed	See analysis
BROADER AI PLATFORM & ENTERPRISE COMPS				
OpenAI	Later Stage VC (S-1 filed Jun 8 '26)	\$852Bn post (Mar 2026); \$180Bn+ raised	\$13.1Bn 2025 rev (Forbes)	~65x trailing
Cohere	Later Stage VC	~\$7Bn (Aug 2025); ~\$1.7Bn raised	\$240M ARR 2025 (CNBC memo)	~29x ARR
Together AI (Nous investor)	Later Stage VC (raising, Mar '26)	\$3.3Bn (Feb 2025); talks at \$7.5Bn pre	~\$1Bn annualized Feb 2026 (Sakra)	~7.5x on ~\$1Bn rev (at \$7.5Bn talks)
STRATEGIC ACQUISITION + PUBLIC COMPS				
MosaicML (Databricks acq.)	Acquired Jul 2023	\$1.3Bn acquisition	~\$20M ARR (CB Insights)	~65x ARR
Run:ai (Nvidia acq.)	Acquired Dec 2024	\$700M acquisition; \$388.8M last mark (Mar '22)	Not disclosed	1.8x vs. stale mark (step-up, not a rev multiple)
CoreWeave (NASDAQ: CRWV)	Public (IPO Mar 2025)	~\$45Bn mkt cap (May '26; moves daily)	\$6.23Bn TTM	~7.3x rev (EV higher w/ debt)

Comp selection anchored to PB's own competitor list for Nous (PBID 572004-37): Prime Intellect, Mistral, Hugging Face, Cohere; all four included. Together AI added because it is both the closest open-model monetization analog and a strategic investor in Nous's Series A (CB Insights). Mistral's ~68x trailing reflects 2025 actuals against a Sep '25 mark; PB carries a Dec '26 estimate of ~\$1.15Bn implying ~12x forward, and the spread is the growth assumption in the price. Hugging Face's mark is 3 years stale; a re-price would likely raise it and lower the multiple. Nous is shown at the reported \$1.5Bn current-round valuation, with the prior \$523M equity and \$1Bn token marks included for context. Multiples in this table mix instruments (equity post-moneys, a token mark, market caps, acquisition prices) and revenue definitions (ARR, run rate, GAAP, estimates); cross-comparisons are directional only. Public market data as of July 2026.

Return and Exit Scenarios

Revenue Multiple Analysis

Prime Intellect is the closest private operating comp at roughly 10x annualized revenue, which would imply approximately \$150M of revenue at Nous's reported \$1.5Bn valuation—likely well above current levels. Other AI comps imply a wider range of roughly \$22M to \$52M, suggesting the round is being priced more on developer adoption, technical credibility, and future monetization than on in-place revenue.

Nous has not disclosed revenue. At an assumed \$40 monthly subscription price, approximately 21,000 paying users would generate \$10M of ARR and 42,000 would generate \$20M. This provides a useful sensitivity against Hermes Agent's large developer base, but it is not confirmed company data. The scenarios below assume entry at \$1.5Bn and gross MOIC before dilution, fees, and expenses.

At a \$1.5Bn entry valuation, returns depend on Nous converting developer adoption into meaningful enterprise revenue and ensuring that equity investors participate in any future token value.

SCENARIO	RETURN / HOLD	OPERATING CASE	EXIT FRAMEWORK	KEY WATCH POINTS
BULL	2.7x-4.0x 3-4 years	Hermes community converts to real Agent / enterprise revenue; Agent holds #1 usage; Psyche network launch adds a second value layer, assuming equity investors participate in the economics.	IPO-track or late-round mark at \$4Bn-\$6Bn	Agent revenue reaches \$50M+, with strong growth and enterprise/API expansion
BASE	1.3x-2.0x 3-5 years	Holds leading independent open-source lab position; Agent revenue grows but stays subscale vs. capital-intensive peers.	Premium strategic acquisition above the MosaicML precedent at \$2Bn-\$3Bn	Agent revenue reaches \$20M-\$50M, but remains subscale relative to larger peers
BEAR	0.3x-0.5x Downside	Frontier free releases erase Hermes differentiation; Agent monetization stalls; token structure complicates equity recovery.	Down round or talent absorption at \$400M-\$750M	Revenue remains below \$20M or stalls, while community adoption weakens

Sources & References

Research date July 2026. PitchBook (PBID 572004-37) Priority 1. Web sources verified Jul '26. Sourcing TBD; no confirmed access path. Full URLs in the source screener.

- 1. PitchBook (PBID 572004-37).** Founded 2022 per PB (press reports 2023 incorporation from a 2022 collective); HQ Austin TX per PB (press cites NYC). Latest round Apr 25 2025 at \$473M pre / \$523M post on \$50M (9.6% acquired). 20 employees (Apr '25). Total raised \$70M incl. \$4.5M debt; investor ownership 30.5%. PB competitor list: Prime Intellect, Mistral AI, Hugging Face, Cohere.
- 2. TechCrunch (Jul 13 2026), rumored round.** Nous finalizing \$75M+ at \$1.5Bn led by Robot Ventures, USV participating (three sources; not yet closed). Founders: Quesnelle, Malhotra, Teknium, Mitra. Hermes Agent \$20-200/mo tiers; ~214K GitHub stars, ~40K forks.
- 3. Fortune (Apr 25 2025), Series A structure.** Paradigm's \$50M priced at a \$1Bn TOKEN valuation, not equity. Prior seed rounds totaled ~\$20M (Distributed Global, North Island, Delphi). ~20-person team at the time. No token issued; unofficial NOUS tokens on Solana DEXs are not company paper (Gate, Apr '26).
- 4. The Block (Apr 25 2025), Series A detail + seed history.** \$50M at \$523M post. Initial \$5.2M seed (Jan '24) led by Distributed Global + OSS Capital with Balaji Srinivasan; unannounced +\$15M Jun '24 round (Delphi, North Island, Vipul Ved Prakash, Raj Gokal). Hermes 3: 50M+ downloads.
- 5. Tech Times / MarkTechPost (May 2026), usage anchor.** Hermes Agent overtook OpenClaw May 10 for #1 on OpenRouter daily rankings (224Bn vs. 186Bn tokens/day); #1 all-time by May 21 at 458Bn/day. Steinberger joined OpenAI Feb '26; OpenClaw moved to an independent foundation with OpenAI as sponsor. NOT an acquisition.
- 6. Hermes 4 Technical Report (arXiv 2508.18255) + VentureBeat.** 405B is a post-train of Llama 3.1 405B (14B from Qwen3). Paper figures: 96.2% MATH-500 / 81.9% AIME'24 reasoning mode; 73.8% / 11.4% standard mode; 57.1% RefusalBench. 405B SFT = 71,616 B200 GPU-hours. Dataset primarily newly synthesized, ~5M samples / 19Bn tokens (no 90% figure in the paper). VentureBeat's 96.3% rounds the paper's 96.2%; Artificial Analysis ranked it #22 of 37 in default mode.
- 7. Nous / Tech Times (Jun-Jul 2026), MoA 2.0.** Mixture of Agents 2.0 shipped in Hermes Agent v0.17.0 (Jun 19) and refined in v0.18.0 (Jul 1). Company claims +8% vs. Claude Opus 4.8 and +11% vs. GPT-5.5 on HermesBench (internal, unpublished; third-party test found ~80x cost per query).
- 8. Prime Intellect Series A, direct peer comp.** \$130M Series A at \$1.0Bn led by Radical Ventures (Jul 8 2026), with NVIDIA Ventures, Intel Capital, Dell Tech Capital, Iconiq. ~\$100M annualized run rate per TechCrunch (company-reported). Priced days before this screen.
- 9. Verification flags from intern draft (corrected in this screen).** (a) OpenClaw was NOT acquired by OpenAI; founder hire plus foundation sponsorship only. (b) A Dillon Roinick IS acknowledged as a Nous colleague in the Hermes 4 technical report, but the 'COO' title and bio could not be verified in press or PB; held out of Management pending confirmation. (c) May '24 '\$210M post' mark unsupported; PB shows \$22M post on the seed. (d) Headcount of 48 unverified (PB: 20). (e) Founded-year and HQ conflicts noted in item 1. (f) Bear-case return math corrected to 0.3-0.5x (\$400M-\$750M exits from \$1.5Bn); all scenarios gross MOIC pre-dilution. (g) External red-team review of this screen incorporated Jul '26.
- 10. Kilo / community sentiment analysis (May 2026).** 1,300+ Reddit comments across 25 top r/openclaw threads: ~35% OpenClaw / ~30% Hermes / ~20% both / ~15% distrust Hermes (astroturfing suspicion). Top criticism: unreliable self-evaluation (+107 upvotes). Security: OpenClaw 9 CVEs in 4 days Mar '26 incl. CVSS 9.9, Koi audit 341/2,857 malicious ClawHub skills; Hermes zero agent CVEs but far less exposure.
- 11. Revenue estimate check (Jul 2026).** No third-party revenue estimate exists: PB blank, ZoomInfo / Growjo no figure, no Sacra note, press explicitly notes no disclosed revenue or subscriber counts (AI Weekly). Business model: Hermes Agent hosted tiers \$20-200/mo plus Nous Portal (300+ models, tool gateway, markup on credits; xAI Grok integration May 15 '26). Founder verification: Malhotra ex-Stanford Brain Stimulation Lab (AI Wiki); Teknium ex-Stability AI, GPTeacher dataset; Mitra founding CEO to President, Prometheia LLC (shivanimitra.com); Quesnelle YaRN cited by 109 papers, used by Meta and DeepSeek (Radical Data Science, Apr '25).
- 12. Comp set sources (Jul 2026).** Hugging Face: \$4.5Bn Aug '23 (stale), \$70M 2023 rev per Sacra, ~\$130M 2024 est. per GetLatka. Cohere: ~\$78n Aug '25 (Radical + Inovia led, ~\$1.7Bn raised); \$240M ARR 2025 per Feb '26 investor memo (CNBC), 70% GM, beat \$200M target. Together AI: \$3.3Bn Feb '25 (GC-led); ~\$1Bn annualized rev Feb '26 (Sacra); in talks at \$7.5Bn pre (Mar '26); invested in Nous's Series A per CB Insights. Cerebras dropped as a comp (chip maker). PB competitor list for Nous: Prime Intellect, Mistral, Hugging Face, Cohere.