

DEAL SCREEN



The AI Native Cloud: open-model inference, training, and GPU infrastructure at scale

San Francisco, CA | Founded 2022 | Together Computer, Inc.

Opportunity	Secondary purchase of Series B-2 Preferred; not public. Seller, size, and transfer mechanics TBC.
Pricing	At the Series C price: \$7.5Bn fully diluted pre-money*. B-2 was struck at the same \$7.5Bn pre in February 2026.
Reference Round	Series C: \$800M closed at \$7.5Bn pre / \$8.3Bn post. Closed March 5, 2026; announced July 1, 2026.
Lead Investors	Aramco Ventures / Prosperity7 (lead); Vista Equity Partners (NEW), NVIDIA, General Catalyst (follow-on), Emergence, March Capital, S Ventures, Pegatron
Sourcing	Cadenza Capital
Step-Up	2.5x from the \$3.3Bn Series B post (Feb 2025)
Total Raised	\$1.34Bn to date; investor ownership 54.3%

** B-2 carries 1x non-participating preference pari passu with existing preferred, 8% non-cumulative dividends, and standard broad-based anti-dilution in the executed term sheet. A secondary transfer is subject to ROFR and co-sale provisions in the existing agreements; confirm before wiring.*

July 2026

STRICTLY CONFIDENTIAL

WHAT THE COMPANY DOES

Together AI is an AI-native cloud built on an AI systems research lab. The company sells the full stack for building and running AI on open and custom models: serverless and dedicated inference APIs, fine-tuning, evals, reserved and instant GPU clusters, and owned "AI factory" data-center capacity. Research is the product engine: the team originated FlashAttention and ships inference optimizations (ATLAS adaptive speculative decoding, tQuant, together.compile) that it claims deliver best-in-class unit economics.

Forge / GPUaaS: Reserved and dedicated GPU capacity for training and inference, now 85% of contracted revenue. The GB300 illustrative build models 80% year-one gross margin and a 42% unlevered-to-levered IRR at 70% LTV debt, with years 4 to 6 upside from re-monetizing depreciated fleet via API inference.

API Inference: Serverless and dedicated endpoints for open models. API inference contracted ARR went \$0 to \$98.6M in FY26, 10x per year since launch, with Cursor as a marquee 1T+ tokens/day customer.

Customers: AI natives on both sides: model builders (Cohere, Pika, Cartesia, ElevenLabs, Suno) and application builders (Cursor, Zoom, Quora, Upwork, Uber, Delivery Hero). Approximately 75 customers as of Q3 FY26; 1M+ developers on platform; 2T inference tokens served daily.

Monetization honesty: "contracted ARR" here is contracted annualized revenue (ACR), not recognized revenue. FY26 recognized revenue was \$144M against \$606M contracted ACR because buildouts convert to revenue as capacity deploys. Both figures are company-provided. Contracted ACR is the metric this cohort is priced on; recognized revenue is the one that has to catch up.

WHY CONSIDER

Cheapest mark in a repricing cohort: The C closed at \$7.5Bn pre in March and announced July 1. Since then Baseten took \$13Bn (Jun '26, ~\$600M reported ARR) and Fireworks printed \$17.5Bn (Jul '26, ~\$800M reported ARR), both roughly 22x. Together at \$8.3Bn on \$1.07Bn contracted ACR is ~7.8x, the cheapest of the private inference cohort. A secondary at the March price in July is buying a stale mark in a market that has moved up.

Revenue at real scale, compounding: Contracted ACR: \$87M FY25 to \$606M FY26 (7x) to \$1.07Bn current; Bookings have since been reported above \$1.15Bn. Base-case model: \$1.8Bn contracted ARR and \$723M recognized revenue in FY27. This is one of the fastest \$0-to-\$1Bn bookings ramps on record, all company-provided.

A genuine research moat in a commodity-risk category: Dao (FlashAttention), Re, Liang, and Zhang give Together the strongest systems-research bench of any neocloud. The efficiency stack is why it can serve inference cheaper than rivals while owning infrastructure, and why the model's GPUaaS margin path (19% today to 40% target) is at least credible.

Syndicate quality and a stated IPO path: Aramco/Prosperity7 lead with NVIDIA, Vista (first entry, T1 growth), and GC re-upping. P7's MD says publicly the company is "headed towards the public markets"; the B-2 documents carry QPO mandatory-conversion mechanics. CoreWeave's IPO (ID8 exit) is the direct precedent path.

Structural downside protection at the entry point: The B-2 security carries a 1x non-participating preference pari passu with all existing preferred, plus 8% non-cumulative dividends and broad-based anti-dilution. At an \$8.3Bn entry against \$1.34Bn of total preference, the paper recovers most capital in mid-case downside scenarios. Combined with the stale-mark pricing, the risk-reward skews better than primary exposure at current cohort marks.

MANAGEMENT TEAM

Vipul Ved Prakash Co-Founder & CEO

Repeat founder (Cloudmark, Topsy, acquired by Apple). Public voice of the open-source AI thesis; also a strategic angel across the ecosystem, including Nous Research.

Ce Zhang Co-Founder & CTO

Professor-track systems researcher (ETH Zurich, University of Chicago). Leads platform engineering and research productization.

Tri Dao Founder & Chief Scientist

Author of FlashAttention, the attention algorithm underlying modern long-context LLM training. Princeton professor. The single largest technical-credibility asset on the roster.

Chris Re and Percy Liang Founders, Professors

Stanford AI faculty (Re: MacArthur Fellow, DAWN/Snorkel lineage; Liang: HELM, foundation-model evaluation). Anchor the research lab identity.

Kai Mak / Charles Zedlewski / Mei Shi CRO / CPO / SVP Finance

Commercial and finance bench scaled with the business; headcount has gone from 282 to 414 in recent months, roughly 130 net adds. Fast hiring is deliberate; integration risk is real.

CAP TABLE AND ROUND HISTORY

Seed (Sep 2022) – \$4.0M @ \$20M post

Lux Capital among earliest backers

Seed / Series A (2023) – ~\$122M across two rounds @ ~\$0.6Bn (Nov '23)

Kleiner Perkins (lead, entered)

T1

Lux Capital (follow-on)

4th Round (Mar 2024) – \$106M @ ~\$1.3Bn

Salesforce Ventures (lead, entered)
Coatue, Emergence (entered)

Series B (Feb 2025) – \$305M @ \$3.3Bn post

General Catalyst (lead, entered)

T1

Prosperity7 (co-lead, entered)

NVIDIA (entered); Salesforce Ventures, Kleiner Perkins, Lux follow-on

Series B-2 (Feb 2026) – \$110M initial @ \$7.5Bn pre; up to \$1Bn additional authorized

Prosperity7 (lead, follow-on)

1x non-participating preference pari passu; 8% non-cumulative dividends

Series C (Mar 2026; announced Jul 2026) – \$800M @ \$7.5Bn pre / \$8.3Bn post

Aramco Ventures / Prosperity7 (lead)

Vista Equity Partners (NEW)

T1

NVIDIA; General Catalyst (follow-on); Emergence; March Capital; S Ventures; Pegatron

FINANCIAL SNAPSHOT

METRIC	FY25 (Jan '25)	FY26 (Jan '26)	CURRENT / FY27E
Contracted ACR	\$87M	\$606M (7.0x YoY)	\$1.07Bn current; \$1.15Bn+ bookings reported
Recognized revenue	\$77M	\$144M	FY27E \$723M (base case)
Adj. EBITDA	n/a	-\$146M	FY27E -\$86M; FY28E +\$403M
Gross margin by product	n/a	GPUaaS 19% / API Ded. 45% / Serverless -172%	Targets: 40% / 50% / 0%
Net revenue retention	0.51 NRR / 0.41 GRR (Q4)	1.07 NRR / 0.91 GRR (Q4)	FY26 compounded: NRR ~1.33, GRR ~0.50
Customers / concentration	n/a	~75; top 5 = ~65% of ACR	#1 customer \$124M, landed in one quarter
Cash / burn	n/a	\$304M cash (Q4); \$13.9M avg monthly burn	Plus \$800M Series C proceeds
Pipeline	n/a	\$5.0Bn potential ACR; \$18.2Bn TCV; \$1.5Bn weighted	One Discovery deal = \$2.1Bn ACR / \$10.7Bn TCV
Valuation mark	\$3.3Bn (Feb '25)	\$7.5Bn pre B-2	\$8.3Bn post Series C (closed)
Headcount	n/a	282	414 (Jul '26)

All operating figures are company-provided, FYE Jan 31. Source attribution and data conflicts are itemized in Sources & References.

DEAL SCORING

CATEGORY	SCORE	RATIONALE
Lead / Round Dynamics	3 / 4	Closed C led by Aramco/P7 with NVIDIA and Vista (NEW T1 growth). Secondary access post-close at round price; not a first-money seat, but a known mark.
Founder / Team Quality	4 / 4	Dao, Re, Liang, Zhang plus a repeat-founder CEO. Best research bench in the category.
Fundamentals	2 / 4	Hypergrowth, but 19% core gross margin, negative EBITDA through FY27, 65% top-5 concentration, and gross retention compounding near 0.50 annually.
Return Potential	3 / 4	~7.8x contracted ACR vs. ~22x at Baseten/Fireworks; stale-mark entry. Offset by 57.6x recognized revenue and capex-driven dilution ahead.
AI Score	4 / 4	FlashAttention lineage, ATLAS, kernels, compiler. The efficiency stack is a real, compounding technical asset.
Terms	2 / 4	B-2 paper is fine (1x pref pari passu). But secondary mechanics unconfirmed: ROFR/co-sale can block or delay, fee load TBC, and no primary information rights transfer automatically.
OVERALL	3.0 / 4.0	(pending diligence)

RISKS & MITIGANTS

Customer concentration: Top 5 customers are ~65% of contracted ACR and the #1 (\$124M) landed in a single quarter. GPU contracts run 1 to 3 years fixed. Mitigant: pipeline is 113 logos and top-5 share should dilute as it converts, but the pipeline itself is concentrated (one deal is 59% of TCV). Weak until proven by renewals.

Margin structure and capex intensity: Core GPUaaS runs 19% gross margin today against a 40% target; serverless is -172%; EBITDA is negative through FY27 and the company plans a 50x infrastructure expansion in five years. That is a continuous-financing model with GPU-depreciation risk; public neocloud peers show depreciation consuming roughly half of revenue. Mitigant: the illustrative GB300 unit economics (80% year-one GM, 42% IRR) are strong if utilization and pricing hold; that is an if.

Retention history: Quarterly gross retention ran 0.74, 0.86, 0.88, 0.91 through FY26, which compounds to roughly 0.50: the existing revenue base loses about half its dollars a year to churn and contraction, and even at the improved Q4 rate the annualized run-rate is only 0.67. Expansion from survivors is what covers it. Quarterly NRR of 1.02, 1.13, 1.09, 1.07 compounds to roughly 1.33, so the existing book still grows about 33% a year on its own. Mitigant: partial. The expansion engine is real and the trend has improved four quarters running. But infrastructure with genuine lock-in does not shed half its base annually; this is a whale book where survivors scale hard and everyone else leaves. The FY26 mega-cohorts replaced \$10M-a-quarter landings with \$200M-a-quarter ones, so the first renewal cycle on those is the test.

Bookings-to-revenue conversion: \$1.07Bn contracted vs. \$144M recognized. The gap converts only as data-center capacity is delivered on time and on cost, in a supply-constrained market. Mitigant: owned "AI factory" strategy reduces dependence on third-party GPU suppliers; it also concentrates execution risk in-house.

Competitive compression: Baseten (\$13Bn) and Fireworks (\$17.5Bn) are freshly funded; Nebius and CoreWeave bring public-scale balance sheets; hyperscalers price aggressively; NVIDIA is invested on every side of the category, including reportedly licensing Groq's LPU stack. Per-token prices deflate structurally. Mitigant: weak. The only durable defense is serving tokens cheaper than rivals, which rests entirely on the research stack and owned infrastructure. Open weights, near-zero switching costs, and a leaky retention base mean there is no second line of defense if that cost advantage closes.

CUSTOMER, RETENTION & PIPELINE DETAIL

Cohort behavior: The Oct '25 cohort (19 customers, \$198.5M at landing) grew to \$218.9M in one quarter; the Jan '26 cohort landed at \$232.9M across 20 customers. Older cohorts expand 1.5x to 3x over 4 to 5 quarters. Landing sizes have step-changed: the business shifted from \$10M-a-quarter cohorts to \$200M-a-quarter cohorts in mid-FY26.

Mix: Forge (GPUaaS) is 84.9% of contracted revenue; API is 15.1% but growing faster (10x per year) at structurally better margins (45% dedicated today vs. 19% GPUaaS). The long-term margin thesis is the mix shift plus fleet re-monetization.

Pipeline quality: \$5.0Bn potential ACR across 113 logos, \$1.5Bn weighted at company win rates. Honest read: a single Discovery-stage opportunity is \$2.1Bn ACR / \$10.7Bn TCV, 43% of pipeline ACR at a 25% stage probability. Strip it and weighted pipeline is roughly \$1.0Bn, still healthy, not the headline.

COMP SET

COMPANY	STAGE	VALUATION	REVENUE	MULTIPLE
CLOSEST PRIVATE COMPS - AI INFERENCE CLOUDS				
Fireworks AI	Later VC	\$17.5Bn (Jul '26); NVIDIA-backed	~\$800M ARR reported	~22x reported
Baseten	Later VC	\$13Bn (Jun '26; \$1.5Bn raised)	~\$600M ARR reported	~22x reported
fal	Later VC	\$4.5Bn	~\$400M ARR reported	~11x reported
Together AI (ID8 entry)	Series C closed	\$8.3Bn post (closed Mar '26; announced Jul '26)	\$1.07Bn contracted ACR / \$144M FY26 recognized	~7.8x ACR; see analysis
PUBLIC NEOCLOUDS				
CoreWeave (CRWV, ID8 exit)	Public (IPO Mar '25)	~\$40Bn to \$48Bn mkt cap (Jul '26, moves daily); EV ~\$80Bn w/ debt	\$6.23Bn TTM	~7x equity / ~13x EV
Nebius (NBIS)	Public	Mkt cap moves daily; not pinned here	~\$1.9Bn AI ARR reported	directional only
STRATEGIC AND SILICON ADJACENCY				
Groq (ID8 portfolio)	Private	NVIDIA LPU license reported at ~\$20Bn (Q2 '26, unverified)	Not disclosed	n/a
Cerebras (CBRS)	Public (IPO '26)	~\$66Bn day-one reported	Not pinned here	n/a

Comp selection: the private cohort is the direct competitive set (Baseten, Fireworks) plus fal for the generative-media wedge; public neoclouds anchor the infrastructure floor; silicon rows are strategic reference only. Multiples in this table mix instruments (private marks, market caps, a license price) and revenue definitions (contracted ACR, reported ARR, TTM GAAP); cross-comparisons are directional only. Private ARR figures are reported, not audited. Data as of July 17, 2026.

RETURN AND EXIT SCENARIOS

Revenue Multiple Analysis. At \$8.3Bn the entry is ~7.8x contracted ACR, 57.6x FY26 recognized revenue, and 11.5x FY27E recognized revenue on the company's base case. The cohort clears at ~22x reported ARR (Baseten, Fireworks), so on the metric the market is actually using, Together is the cheap name; on recognized GAAP-style revenue it is the expensive one, and the spread is exactly the bookings-to-revenue conversion risk. The seller-side materials model bear/base/mgmt/bull exits of \$12.45Bn / \$24.9Bn / \$49.8Bn / \$83Bn; those are seller-side figures and we do not adopt them. ID8's own read below assumes entry at \$8.3Bn post-equivalent, gross MOIC before fees, carry, and future dilution, and notes that B-2's 1x pari passu preference provides real, not theoretical, downside cushion at this entry point.

SCENARIO	RETURN / HOLD	OPERATING CASE	EXIT FRAMEWORK	KEY WATCH POINTS
BULL	3x to 4x 2 to 4 years	FY28 to FY29 revenue ramps toward \$2Bn to \$3.7Bn per model; GPUaaS margin reaches the 40% target; API mix shift compounds; IPO window opens on the CoreWeave precedent.	IPO or late round at \$25Bn to \$35Bn (8x to 12x forward revenue)	GPUaaS margin trajectory; top-5 concentration falling below 50%; capacity delivered on schedule
BASE	1.8x to 2.5x 3 to 5 years	Growth persists but the margin grind is slow; bookings convert with delays; category multiples compress toward public-neocloud reality.	IPO or strategic at \$15Bn to \$20Bn	NRR holding above 1.05; renewal outcomes on the FY26 mega-cohorts; financing terms for the 50x buildout
BEAR	0.5x to 1x Downside	Token-price deflation plus a top-5 loss; serverless losses persist; capex refinancing tightens across the neocloud complex.	Down round or consolidation at \$4Bn to \$6Bn; the 1x pref recovers most capital unless the pref stack itself is breached	Any top-5 non-renewal; GPU depreciation outpacing re-monetization; debt markets for GPU collateral

SOURCES & REFERENCES

Research date July 17, 2026. Current-round and access terms per ID8 internal information (Priority 1 for the live deal); PitchBook (PBID 518020-12) Priority 1 for the historical record; dataroom figures company-provided. Full URLs in the source screener.

- ID8 internal (Jul '26).** Access is a secondary purchase of Series B-2 Preferred priced at the Series C price, per ID8 deal discussions. Seller identity, block size, broker, fee load, and ROFR/co-sale posture all TBC before any commitment.
- PitchBook (PBID 518020-12).** Series C: \$800M closed Mar 5, 2026 at \$7.5Bn pre / \$8.3Bn post, led by Prosperity7 and Aramco Ventures, SE Ventures participating (PB note). Prior: \$305M Feb '25 (5th round), \$106M Mar '24 (4th round), \$4M seed Sep '22 @ \$20M est. Total raised \$1.34Bn; investor ownership 54.3%; 414 employees Jul 13 '26; PB revenue mark \$44M TTM Dec '24 (stale, different basis).
- P7 Term Sheet (executed Feb 24, 2026, dataroom).** Series B-2 at \$7.5Bn fully diluted pre; P7 \$110M initial plus up to \$1Bn additional within 90 days; 1x non-participating pref pari passu; 8% non-cumulative dividends; QPO mandatory conversion; board of 9 (P7 seat subject to CFIUS review); ROFR/co-sale per existing agreements.
- Company deck (dataroom).** Contracted ARR \$606M FY26 (6x YoY stated), current \$1.07Bn; \$18Bn+ TCV pipeline; FY27 base case \$1.8Bn; Series C target ~\$800M @ \$7.5Bn pre; 2T tokens/day; 1M+ developers; 282 employees; \$548M raised pre-C.
- Company model 2026-03 (dataroom).** FY26A: ACR \$605.6M, revenue \$144.1M, adj. EBITDA -\$146.4M. FY27E: ACR \$1,844.9M, revenue \$723.1M, EBITDA -\$86.3M. FY28E: revenue \$1,974.9M, EBITDA +\$403.3M. GPUaaS margin 19% to 40% path; utilization loss \$42.5M FY26.
- Retention Waterfall / Cohorted ACR (dataroom).** Q4 FY26: ending ACR \$594.3M (waterfall basis; \$605.9M total incl. self-serve), GRR 0.91, NRR 1.07. FY25 Q4: churn \$53.7M, GRR 0.41, footnoted \$47M one-time. Oct '25 cohort \$198.5M to \$218.9M in one quarter; Jan '26 cohort \$232.9M.
- Top 50 Customers / Customer Cube (dataroom).** Top 5 = \$388M of \$594M (65%); #1 customer \$124.25M (AI model builder, Forge) landed in Q4 FY26 from zero; #2 \$113M public tech enterprise; #3 \$73M AI application on API. Median ACR \$477K vs. average \$6.1M: a whale-driven book.
- Open Pipeline / Stage Definitions (dataroom).** \$5.0Bn potential ACR, \$18.2Bn TCV, \$1.5Bn weighted at stage probabilities (15% to 85%). Single Discovery-stage opportunity: \$2.13Bn ACR / \$10.66Bn TCV, 59% of total TCV at 25% probability.
- GPU ROI illustrative (dataroom).** GB300, 2,500 GPUs, \$4.25/GPU-hr, \$75K capex per GPU (\$187.5M total), 70% LTV debt at 9%: 80% year-one GM, ~76% EBITDA margin, 42% IRR over 6-year life with pricing decay to \$3.64 by year 6.
- Series C press (Jul 1, 2026; BusinessWire, TechCrunch, Verdict, DCD).** \$800M at \$8.3Bn post announced; Aramco Ventures lead; Vista, Emergence, GC, March Capital, NVIDIA, S Ventures, Pegatron participating; bookings above \$1.15Bn last quarter; 50x infrastructure expansion planned; P7 MD: "headed towards the public markets." Mar-vs-Jul close date reconciled to staged closings.
- Comps (Jul '26).** Fireworks \$17.5Bn (CNBC, Jul 16). Baseten \$13Bn on \$1.5Bn raised (press, Jun '26); earlier \$5Bn Feb '26 (Sactra). fal \$4.5Bn / ~\$400M ARR (secondary research). CoreWeave ~\$40Bn to \$48Bn mkt cap, \$6.23Bn TTM revenue, EV ~\$80Bn (CNBC/macrotrends/multiples.vc, dates vary within Jul '26). Nebius ~\$1.9Bn AI ARR reported. Groq/NVIDIA ~\$20Bn LPU license and Cerebras ~\$66Bn IPO are unverified secondary reporting; confirm before quoting externally.
- Verification flags.** (a) Deck says C "target" while PB and press show it closed; screen treats the round as closed. (b) Press names Aramco lead; PB note adds Prosperity7 and SE Ventures; board docs give P7 the seat; treated as one Aramco/P7 lead. (c) Headcount 282 (deck) vs. 414 (PB): both shown. (d) Revenue \$77M FY25 (deck) vs. PB \$44M TTM Dec '24: different calendars and bases; company figure used, departure flagged. (e) Nov '23 ~\$0.6Bn mark is broker-chart only. (f) External red-team review of this screen not yet run.